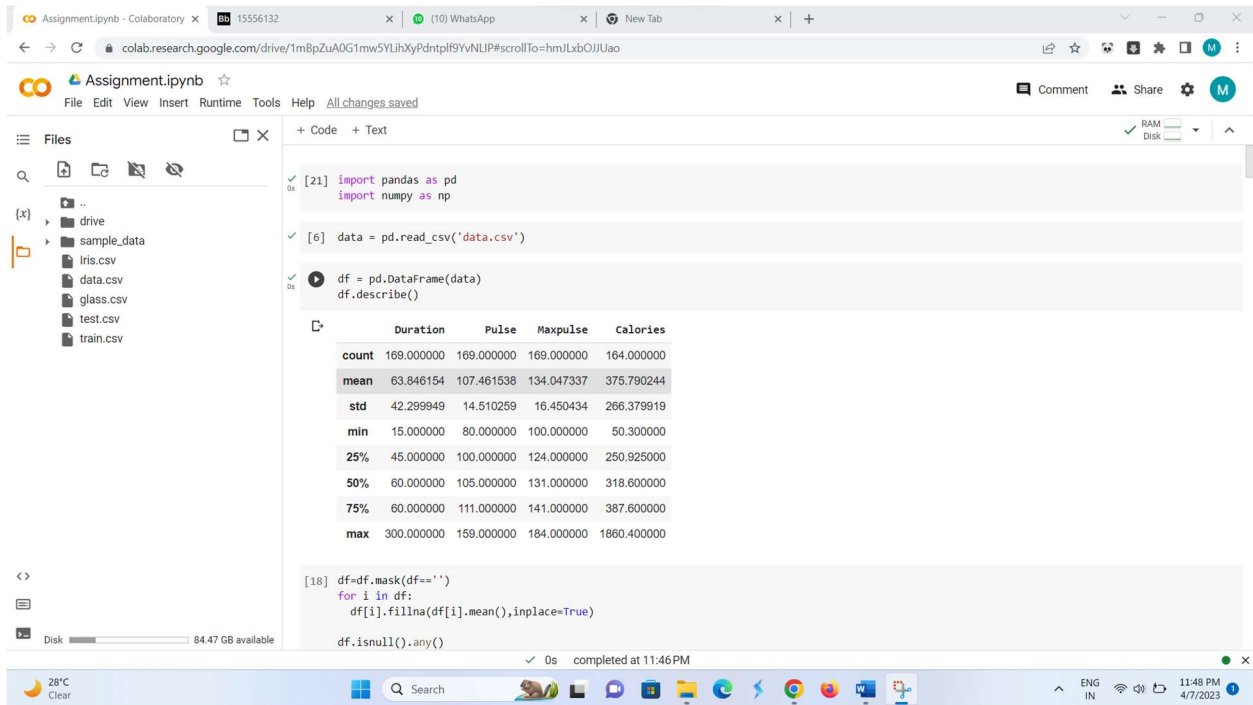


Name: Mythresh Maddina

700: 700741162

Github Link: https://github.com/MythreshM/CS5710_Assignment4

Question1:



The screenshot shows a Google Colab notebook titled "Assignment.ipynb". The left sidebar displays a file explorer with a folder named "drive" containing a subfolder "sample_data" and several CSV files: "iris.csv", "data.csv", "glass.csv", "test.csv", and "train.csv". The main code area contains three cells:

- Cell [21]: `import pandas as pd`
`import numpy as np`
- Cell [6]: `data = pd.read_csv('data.csv')`
- Cell [0]: `df = pd.DataFrame(data)`
`df.describe()`

The output of the third cell is a summary table for the DataFrame:

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

Below the table, the notebook shows the execution of a loop to fill missing values:

```
[18] df=df.mask(df=='')
for i in df:
    df[i].fillna(df[i].mean(),inplace=True)
df.isnull().any()
```

The bottom status bar indicates the notebook is completed at 11:46 PM. The Windows taskbar at the very bottom shows the date as 4/7/2023 and the time as 11:48 PM.

Assignment.ipynb - Colaboratory

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNLP#scrollTo=hmLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- drive
- sample_data
 - Iris.csv
 - data.csv
 - glass.csv
 - test.csv
 - train.csv

```
df=df.mask(df=='')
for i in df:
    df[i].fillna(df[i].mean(),inplace=True)

df.isnull().any()

10790
18161
22654
63508,5512195122
Duration False
Pulse False
Maxpulse False
Calories False
dtype: bool
```

[19] #4 Here we select Duration and pulse

```
df.agg({'Duration': ['min', 'max', 'count', 'mean'], 'Calories': ['min', 'max', 'count', 'mean']})
```

	Duration	Calories
min	15.000000	50.300000
max	300.000000	1860.400000
count	169.000000	169.000000
mean	63.846154	375.790244

[16] #5 Filter the dataframe to select the rows with calories values between 500 and 1000

```
df.loc[(df['Calories']>500) & (df['Calories']<1000)]
```

0s completed at 11:46 PM

28°C Clear

Search

ENG IN 11:48 PM 4/7/2023

Assignment.ipynb - Colaboratory

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNLP#scrollTo=hmLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- drive
- sample_data
 - Iris.csv
 - data.csv
 - glass.csv
 - test.csv
 - train.csv

```
mean 63.846154 375.790244
```

[5] #5 Filter the dataframe to select the rows with calories values between 500 and 1000

```
df.loc[(df['Calories']>500) & (df['Calories']<1000)]
```

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
90	180	101	127	600.1
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

0s completed at 11:46 PM

28°C Clear

Search

ENG IN 11:48 PM 4/7/2023

Assignment.ipynb - Colaboratory

15556132 (10) WhatsApp New Tab

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmJLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Files

drive

sample_data

Iris.csv

data.csv

glass.csv

test.csv

train.csv

Code

```
#6 Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
df.loc[(df['Calories']>500)&(df['Pulse']<100)]
```

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

```
[ ] #7 Create a new "df_modified" dataframe except "Maxpulse".
df_modified=df[['Calories','Pulse','Calories']]
df_modified.head()
```

```
[ ] #8 delete max_pulse column in df use del
del df['Maxpulse']
```

completed at 11:46 PM

28°C Clear

Search

ENG IN

11:49 PM 4/7/2023

Assignment.ipynb - Colaboratory

15556132 (10) WhatsApp New Tab

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmJLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Files

drive

sample_data

Iris.csv

data.csv

glass.csv

test.csv

train.csv

Code

```
[ ] #8 delete max_pulse column in df use del
del df['Maxpulse']
```

```
#9 Converting the datatype of Calories column to int datatype using astype
df['Calories'] = df['Calories'].astype(np.int64)
df.dtypes
```

	Duration	Pulse	Maxpulse	Calories	dtype
	int64	int64	int64	int64	object

```
[28] #10 Using pandas create a scatter plot for the two columns (Duration and Calories)
df.plot.scatter(x='Duration', y='Calories',c='blue')

<Axes: xlabel='Duration', ylabel='Calories'>
```

completed at 11:46 PM

28°C Clear

Search

ENG IN

11:49 PM 4/7/2023

Question2:

Assignment.ipynb - Colaboratory

15556132

(10) WhatsApp

New Tab

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmJLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Files

drive

sample_data

Iris.csv

data.csv

glass.csv

test.csv

train.csv

Question 2:

1. Titanic Dataset

```
[58] import pandas as pd
from sklearn import preprocessing
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, recall_score, precision_score, classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")

[30] df=pd.read_csv("train.csv")
df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

0s completed at 11:46 PM

28°C Clear

Search

ENG IN

11:49 PM 4/7/2023

Assignment.ipynb - Colaboratory

15556132

(10) WhatsApp

New Tab

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmJLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Files

drive

sample_data

Iris.csv

data.csv

glass.csv

test.csv

train.csv

```
# converting categorical data to numerical values for calculating correlation
label_encoder = preprocessing.LabelEncoder()
df['Sex'] = label_encoder.fit_transform(df.Sex.values)

#calculating correlation for 'Survived' and 'Sex' in given data
correlation_value= df['Survived'].corr(df['Sex'])
print(correlation_value)

-0.5433513806577555

A) Yes, we should keep survived and sex features which helps to classify data accurately.

[33] # Display Correlation
correlation_matrix = df.corr()
print(correlation_matrix)
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651

	Parch	Fare
PassengerId	-0.001652	0.012658
Survived	0.081629	0.257307

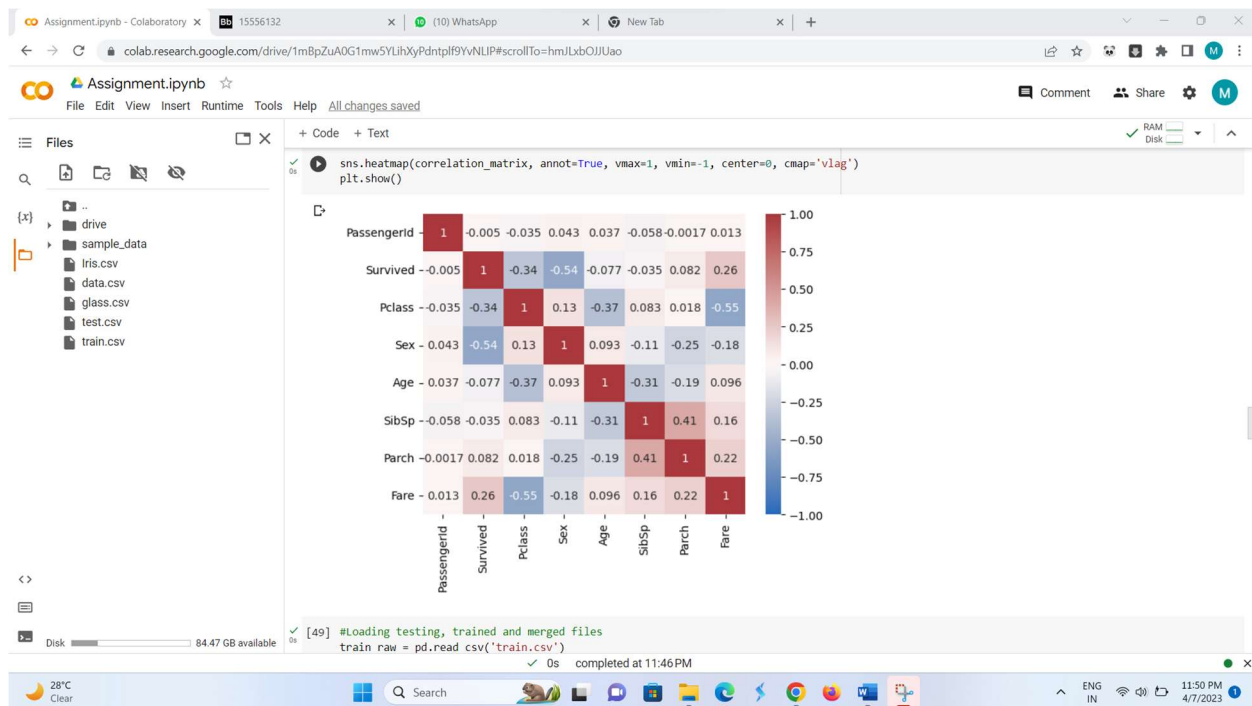
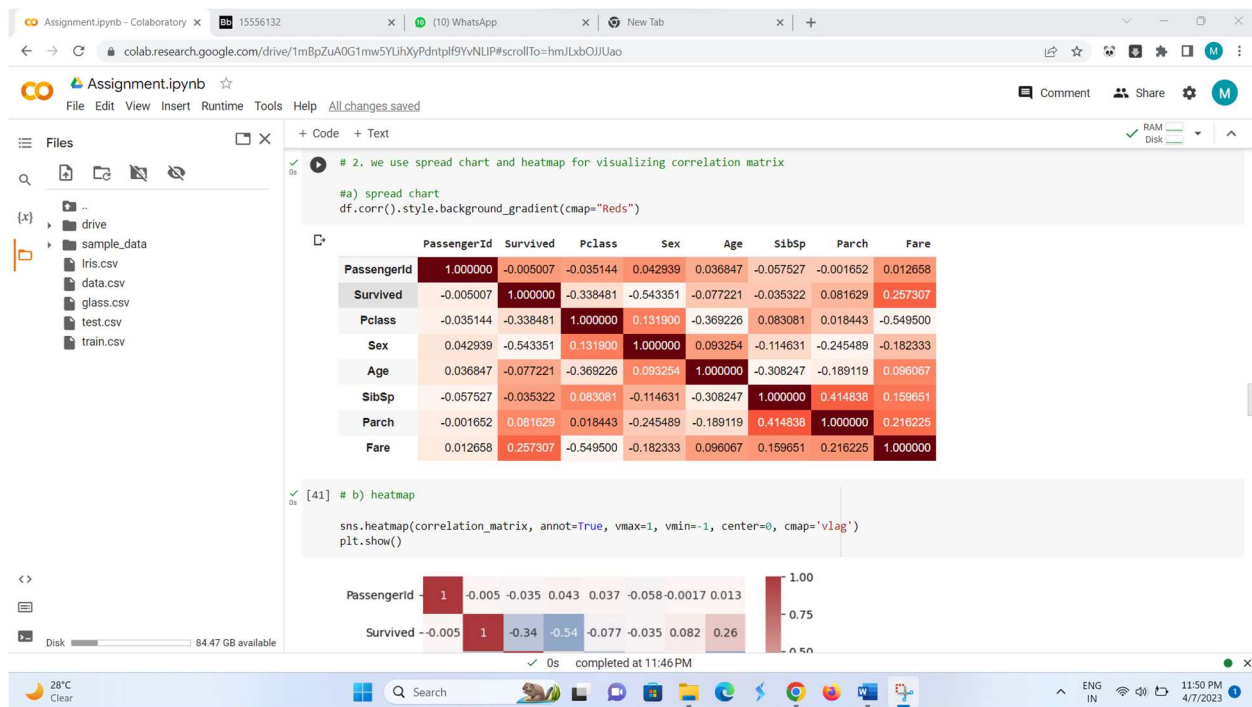
0s completed at 11:46 PM

28°C Clear

Search

ENG IN

11:50 PM 4/7/2023



The screenshot displays a Google Colaboratory notebook titled "Assignment.ipynb". The interface includes a file explorer on the left showing a directory structure with files like "Iris.csv", "data.csv", "glass.csv", "test.csv", and "train.csv". The main area contains four code cells:

- Cell 0:** Loading training and testing files, merging them into a single DataFrame, and defining features and target variables. It includes data cleaning steps like replacing "female" with "male" and "S" with "C" in the "Embarked" column.
- Cell 50:** Dropping missing values from the training set using `train.dropna` and `train.drop`.
- Cell 51:** Splitting the training data into training and validation sets using `train_test_split`.
- Cell 52:** Initializing a Gaussian Naive Bayes classifier with `GaussianNB()`.

The notebook status bar at the bottom indicates it was completed at 11:46 PM. The system tray at the very bottom shows a temperature of 28°C, a search bar, and the date/time 11:50 PM on 4/7/2023.

Question 3:

Assignment.ipynb - Colaboratory

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmJLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- drive
- sample_data
 - Iris.csv
 - data.csv
 - glass.csv
 - test.csv
 - train.csv

```
y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score

print('accuracy is', accuracy_score(Y_val, y_pred))
```

	precision	recall	f1-score	support
0.0	0.79	0.80	0.80	85
1.0	0.70	0.69	0.70	58
accuracy			0.76	143
macro avg	0.75	0.74	0.75	143
weighted avg	0.75	0.76	0.75	143

```
[[68 17]
 [18 40]]
accuracy is 0.7552447552447552
```

Glass Dataset

```
glass=pd.read_csv("glass.csv")
glass.head()
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1

0s completed at 11:46 PM

28°C Clear

Search

ENG IN

11:51 PM 4/7/2023

Assignment.ipynb - Colaboratory

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmJLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- drive
- sample_data
 - Iris.csv
 - data.csv
 - glass.csv
 - test.csv
 - train.csv

```
glass.corr().style.background_gradient(cmap="Reds")
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
RI	1.000000	-0.191885	-0.122274	-0.407326	-0.542052	-0.289833	0.810403	-0.000386	0.143010	-0.164237
Na	-0.191885	1.000000	-0.273732	0.156794	-0.069809	-0.266087	-0.275442	0.326603	-0.241346	0.502898
Mg	-0.122274	-0.273732	1.000000	-0.481799	-0.165927	0.005396	-0.443750	-0.492262	0.083060	-0.744993
Al	-0.407326	0.156794	-0.481799	1.000000	-0.005524	0.325958	-0.259592	0.479404	-0.074402	0.598829
Si	-0.542052	-0.069809	-0.165927	-0.005524	1.000000	-0.193331	-0.208732	-0.102151	-0.094201	0.151565
K	-0.289833	-0.266087	0.005396	0.325958	-0.193331	1.000000	-0.317836	-0.042618	-0.007719	-0.010054
Ca	0.810403	-0.275442	-0.443750	-0.259592	-0.208732	-0.317836	1.000000	-0.112841	0.124968	0.000952
Ba	-0.000386	0.326603	-0.492262	0.479404	-0.102151	-0.042618	-0.112841	1.000000	-0.058692	0.575161
Fe	0.143010	-0.241346	0.083060	-0.074402	-0.094201	-0.007719	0.124968	-0.058692	1.000000	-0.188278
Type	-0.164237	0.502898	-0.744993	0.598829	0.151565	-0.010054	0.000952	0.575161	-0.188278	1.000000

```
sns.heatmap(correlation_matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```

PassengerId 1 -0.005 -0.035 0.043 0.037 -0.058 -0.0017 0.013

Survived -0.005 1 -0.34 -0.54 -0.077 -0.035 0.082 0.26

Pclass -0.035 -0.34 1 0.13 -0.37 0.083 0.018 -0.55

0s completed at 11:46 PM

28°C Clear

Search

ENG IN

11:51 PM 4/7/2023

Assignment.ipynb - Colaboratory

15556132

(10) WhatsApp

New Tab

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Files

drive

sample_data

Iris.csv

data.csv

glass.csv

test.csv

train.csv

Code

```
features = ['Rl', 'Na', 'Mg', 'Al', 'Si', 'K', 'Ca', 'Ba', 'Fe']
target = 'Type'
X_train, X_val, Y_train, Y_val = train_test_split(glass[features], glass[target], test_size=0.2, random_state=1)
classifier = GaussianNB()
classifier.fit(X_train, Y_train)
y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
print('accuracy is', accuracy_score(Y_val, y_pred))
```

	precision	recall	f1-score	support
1	0.90	0.95	0.92	19
2	0.92	0.92	0.92	12
3	1.00	0.50	0.67	6
5	0.00	0.00	0.00	1
6	1.00	1.00	1.00	1
7	0.75	0.75	0.75	4
accuracy			0.84	43
macro avg	0.76	0.69	0.71	43
weighted avg	0.89	0.84	0.85	43

```
[[18 1 0 0 0 0]
 [ 1 11 0 0 0 0]
 [ 1 0 3 2 0 0]
 [ 0 0 0 0 1 1]
 [ 0 0 0 0 1 0]
 [ 0 0 0 1 0 3]]
```

0s completed at 11:46 PM

28°C Clear

Search

ENG IN

11:52 PM 4/7/2023

Assignment.ipynb - Colaboratory

15556132

(10) WhatsApp

New Tab

colab.research.google.com/drive/1mBpZuA0G1mw5YLhXyPdntplf9YvNUP#scrollTo=hmLxbOJJUao

Assignment.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

Files

drive

sample_data

Iris.csv

data.csv

glass.csv

test.csv

train.csv

Code

```
from sklearn.svm import SVC, LinearSVC

classifier = LinearSVC()
classifier.fit(X_train, Y_train)
y_pred = classifier.predict(X_val)
# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(Y_val, y_pred))
```

	precision	recall	f1-score	support
1	0.82	0.95	0.88	19
2	0.70	0.58	0.64	12
3	0.00	0.00	0.00	6
5	0.00	0.00	0.00	1
6	0.14	1.00	0.25	1
7	0.00	0.00	0.00	4
accuracy			0.60	43
macro avg	0.28	0.42	0.29	43
weighted avg	0.56	0.60	0.57	43

```
[[18 1 0 0 0 0]
 [ 4 7 0 0 1 0]
 [ 0 2 0 3 1 0]
 [ 0 0 0 0 1 0]
 [ 0 0 0 0 1 0]
 [ 0 0 0 1 3 0]]
```

accuracy is 0.6046511627906976

0s completed at 11:46 PM

28°C Clear

Search

ENG IN

11:52 PM 4/7/2023